

# Error Decomposition and Boosting

Wei Wang @ CSE, UNSW

April 11, 2020

# Overview

- Error decomposition
  - Bias-variance (+noise) decomposition
  - Estimation-approximation decomposition
- Adaboost:
  - We derive the Adaboosting algorithm (for classification) from the Gradient Boosting perspective.


# Parameter Estimation

- By now, you should realize that we are essentially doing parameter estimation in Machine Learning. The two inputs are:

- Parameterized function family:  $f(\mathbf{x}; \theta)$
- Training data:  $\mathbf{X}$

The output is  $\hat{\theta}$ .

- Examples:
  - MLE
  - MAP
- How to evaluate the quality of the estimator  $\hat{\theta}$ ?
  - $\hat{\theta}$  is a random variable wrt.  $\mathbf{X}$ !
  - Bias and Variance, i.e.,  $\mathbf{E}_{\mathbf{X}} [\hat{\theta}]$  and  $\mathbf{Var}_{\mathbf{X}} [\hat{\theta}]$



# Bias-Variance Tradeoff (on the Parameter)

- Estimated difference between a specific estimate and the true parameter:

- Change of notation:  $D$  for  $\mathbf{X}$ ,  $\hat{\theta} \stackrel{\text{def}}{=} h(D)$
- Blue symbols are random variables.

$$\begin{aligned}
 \mathbf{E}_D \left[ (h(D) - \theta)^2 \right] &= \mathbf{E}_D \left[ (h(D) - \mu + \mu - \theta)^2 \right] \\
 &= \mathbf{E}_D \left[ (h(D) - \mu)^2 \right] + \mathbf{E}_D \left[ (\mu - \theta)^2 \right] + \underbrace{2 \mathbf{E}_D \left[ (h(D) - \mu) \cdot (\mu - \theta) \right]}_{\text{cross term}} \\
 &= \underbrace{\mathbf{E}_D \left[ (h(D) - \mathbf{E}_D \left[ h(D) \right])^2 \right]}_{\text{variance}} + \underbrace{\mathbf{E}_D \left[ (\mathbf{E}_D \left[ h(D) \right] - \theta)^2 \right]}_{\text{bias}}
 \end{aligned}$$

- Elimination of the cross term:

$$\begin{aligned}
 \mathbf{E}_D \left[ (h(D) - \mu) \cdot (\mu - \theta) \right] &= (\mu - \theta) \cdot \mathbf{E}_D \left[ (h(D) - \mu) \right] \\
 &= (\mu - \theta) \cdot (\mathbf{E}_D \left[ h(D) \right] - \mu)
 \end{aligned}$$

- Choose  $\mu \stackrel{\text{def}}{=} \mathbf{E}_D \left[ h(D) \right]$  (not a r.v.), then this term vanishes

# Understanding

$$\mathbf{E}_D \left[ (h(D) - \theta)^2 \right] = \underbrace{\mathbf{E}_D \left[ (h(D) - \mu)^2 \right]}_{\text{variance}} + \underbrace{\mathbf{E}_D \left[ (\mu - \theta)^2 \right]}_{\text{bias}}$$

- $\mu$  is the (“long-term”) expected parameter inferred from training data
- Can you find approximately where  $\mu$ s are in the previous 1D examples?

# Bias-Variance Decomposition for Regression


- Consider the regression with  $\ell_2$  loss:  $\ell = (h - f)^2$ .
  - Note that  $f$  and  $h$  are both functions, but we hide the input  $D$ .
  - Also note that summation is not needed, as the input is  $D$ .
- Consider its expectation.

$$\begin{aligned}\mathbf{E} [\ell] &= \mathbf{E} [(h - f)^2] = \mathbf{E} [((h - \bar{h}) + (\bar{h} - f))^2] \\&= \mathbf{E} [(h - \bar{h})^2] + \mathbf{E} [(\bar{h} - f)^2] + \mathbf{E} [2(h - \bar{h}) \cdot (\bar{h} - f)] \\&= \underbrace{\mathbf{E} [(h - \bar{h})^2]}_{\text{variance}} + \underbrace{\mathbf{E} [(\bar{h} - f)^2]}_{\text{bias}}\end{aligned}$$

- Bias measures how much the prediction (averaged over all training data sets) differs from the desired regression function.
- Variance measures how much the predictions for individual training data sets vary around their average.


# Understanding

- There is a trade-off between bias and variance. As we increase model complexity (from **blue function family** to **green family**)
  - bias decreases (more likely to fit the training data better) and
  - variance increases (the choice of the learned function varies more with training data)
- No such clean decomposition for classification, but the trade-off still holds.




## Example /1

$$\mathcal{H}_0 : y = c \text{ and } \mathcal{H}_1 : y = a \cdot x + b$$



$$\begin{aligned} \mathcal{H}_0 \\ \text{bias} &= 0.50 \\ \text{var} &= 0.25 \\ \hline E_{\text{out}} &= 0.75 \quad \checkmark \end{aligned}$$




$$\begin{aligned} \mathcal{H}_1 \\ \text{bias} &= 0.21 \\ \text{var} &= 1.69 \\ \hline E_{\text{out}} &= 1.90 \end{aligned}$$



## Example /2

2 Data Points



$$\mathcal{H}_0$$

$$\text{bias} = 0.50;$$

$$\text{var} = 0.25.$$

$$E_{\text{out}} = 0.75 \quad \checkmark$$


$$\mathcal{H}_1$$

$$\text{bias} = 0.21;$$

$$\text{var} = 1.69.$$

$$E_{\text{out}} = 1.90$$

5 Data Points



$$\mathcal{H}_0$$

$$\text{bias} = 0.50;$$

$$\text{var} = 0.1.$$

$$E_{\text{out}} = 0.6$$

$$\mathcal{H}_1$$

$$\text{bias} = 0.21;$$

$$\text{var} = 0.21.$$

$$E_{\text{out}} = 0.42 \quad \checkmark$$

# The General Version

- Considering the noise, i.e.,

$$\hat{f}(\mathbf{x}) = f(\mathbf{x}) + \underbrace{\epsilon}_{\text{noise}}$$

then the decomposition can be decomposed as

$$\mathbf{E} \left[ (\hat{f} - f)^2 \right] = \mathbf{Var} \left[ \hat{f} \right] + (\mathbf{E} \left[ \hat{f} \right] - f)^2 + \underbrace{\mathbf{Var} \left[ \epsilon \right]}_{\text{noise}}$$

- The additional noise component represents *irreducible error*.

# Estimation-Approximation Error Decomposition


- Use linear regression as an example.

$$\begin{aligned}
 \mathbf{E} \left[ (\hat{f} - f)^2 \right] &= \mathbf{E} \left[ ((\hat{f} - \hat{f}^*) + (\hat{f}^* - f))^2 \right] \\
 &= \mathbf{E} \left[ (\hat{f} - \hat{f}^*)^2 \right] + \mathbf{E} \left[ (\hat{f}^* - f)^2 \right] + \underbrace{2 \mathbf{E} \left[ (\hat{f} - \hat{f}^*) \cdot (\hat{f}^* - f) \right]}_{\text{cross term}} \\
 &= \underbrace{\mathbf{E} \left[ (\hat{f} - \hat{f}^*)^2 \right]}_{\text{estimation error}} + \underbrace{\mathbf{E} \left[ (\hat{f}^* - f)^2 \right]}_{\text{approximation error}}
 \end{aligned}$$


- The expectation is wrt. the distribution from which training and test data are sampled.
- The cross term must be 0 as the error of the best possible approximation ( $\hat{f}^* - f$ ) is independent of the error of any linear function of the data (hence  $\hat{f} - \hat{f}^*$ ).
- General case

$$\ell_D^{0-1}[h_S] = \underbrace{\left( \ell_D^{0-1}[h_S] - \inf_{h \in \mathcal{H}} \ell_D^{0-1}[h] \right)}_{\text{Estimation Error in } \mathcal{H}} + \underbrace{\left( \inf_{h \in \mathcal{H}} \ell_D^{0-1}[h] - \ell_D^{0-1,*} \right)}_{\text{Approximation error of } \mathcal{H}} + \underbrace{\ell_D^{0-1,*}}_{\text{Irreducible Error}}$$



# Example



$\ell$



$n$



# Additive Model

- $H_t(\mathbf{x}) = \sum_{k=1}^t \alpha_k h_k(\mathbf{x})$
- Each  $h_i$  belong to a function family  $\mathcal{H}$  (typically a weak functional class)
- Intuitively, we judiciously increase the complexity of the function class one step at a time (rather than going towards a very complex function class that will easily get overfitted).


# Naïve (Regression) Method

- Training data:  $(\mathbf{x}_i, y_i)$
- Current model's prediction:  $(\mathbf{x}_i, H_t(\mathbf{x}_i))$
- Learn a function  $\alpha_{t+1} h_{t+1}()$  such that it approximates  $y_i - H_t(\mathbf{x}_i)$  well.
  - A new regression problem with the effective training data:  $(\mathbf{x}_i, y_i - H_t(\mathbf{x}_i))$
- What's the problem with this method?

# Exponential Loss

- 0 – 1 loss is hard to optimize directly.
- Exponential loss upper bounds the 0-1 loss:

$$\ell(H_t) = \sum_{k=1}^n \exp(-y_i H_t(\mathbf{x}_i))$$



- Black: 0-1 Loss
- Blue: Hinge Loss
- Red: Logistic Loss
- Green: Exponential Loss

From  $H_t(\mathbf{x})$  to  $H_{t+1}(\mathbf{x})$ 

$$\begin{aligned}\ell_{t+1} &= \sum_{i=1}^n \exp(-y_i H_{t+1}(\mathbf{x}_i)) \\ &= \sum_{i=1}^n \exp(-y_i [H_t(\mathbf{x}_i) + \alpha_{t+1} h_{t+1}(\mathbf{x}_i)]) \\ &= \sum_{i=1}^n \exp(-y_i H_t(\mathbf{x}_i)) \cdot \exp(-y_i \alpha_{t+1} h_{t+1}(\mathbf{x}_i)) \\ &\stackrel{\text{def}}{=} \sum_{i=1}^n w_t(i) \cdot \exp(-y_i \alpha_{t+1} h_{t+1}(\mathbf{x}_i))\end{aligned}$$

- To minimize the loss  $\ell_{t+1}$ , we have two “variables” to tune:  $h_{t+1}()$  and  $\alpha_{t+1}$  (note that  $w_t(i)$  is a constant by now)



Choose  $h_{t+1}()$  to minimize  $\ell_{t+1}$ 

We will omit the subscripts on  $\alpha$  and  $h()$

$$\begin{aligned}
 \ell_{t+1} &= \sum_{i=1}^n w_t(i) \cdot \exp(-y_i \alpha h(\mathbf{x}_i)) \\
 &= \sum_{i|y_i=h(\mathbf{x}_i)} w_t(i) \cdot e^{-\alpha} + \sum_{i|y_i \neq h(\mathbf{x}_i)} w_t(i) \cdot e^{\alpha} \\
 &= \sum_{i|y_i=h(\mathbf{x}_i)} w_t(i) \cdot e^{-\alpha} + \sum_{i|y_i \neq h(\mathbf{x}_i)} w_t(i) \cdot e^{-\alpha} - \sum_{i|y_i \neq h(\mathbf{x}_i)} w_t(i) \cdot e^{-\alpha} + \sum_{i|y_i \neq h(\mathbf{x}_i)} w_t(i) \cdot e^{\alpha} \\
 &= e^{-\alpha} \cdot \left( \sum_{i|y_i=h(\mathbf{x}_i)} w_t(i) + \sum_{i|y_i \neq h(\mathbf{x}_i)} w_t(i) \right) + (e^{\alpha} - e^{-\alpha}) \cdot \sum_{i|y_i \neq h(\mathbf{x}_i)} w_t(i) \\
 &= e^{-\alpha} \cdot \sum_{i=1}^n w_t(i) + (e^{\alpha} - e^{-\alpha}) \cdot \sum_{i=1}^n w_t(i) \mathbf{1}[y_i \neq h(\mathbf{x}_i)] \\
 &= e^{-\alpha} + (e^{\alpha} - e^{-\alpha}) \cdot \sum_{i=1}^n w_t(i) \mathbf{1}[y_i \neq h(\mathbf{x}_i)] \quad (\dagger)
 \end{aligned}$$

## Minimization /1

$$\ell_{t+1} = e^{-\alpha} + (e^{\alpha} - e^{-\alpha}) \cdot \sum_{i=1}^n w_t(i) \mathbf{1}[y_i \neq h(\mathbf{x}_i)] \quad (\dagger)$$

- Choose  $h_{t+1}$  (within the function family) to minimize  $\epsilon_t \stackrel{\text{def}}{=} \sum_{i=1}^n w_t(i) \mathbf{1}[y_i \neq h(\mathbf{x}_i)]$ .
  - $h_{t+1}()$  minimizes the weighted error.
- After  $h_{t+1}()$  is learned (i.e., fixed), need to determine the best  $\alpha_{t+1}$ :

$$\frac{\partial \ell_{t+1}}{\partial \alpha} = e^{\alpha} \epsilon_t + e^{-\alpha} \epsilon_t - e^{-\alpha}$$

Setting it to 0 give us:

$$\alpha_{t+1} = \frac{1}{2} \ln \frac{1 - \epsilon_t}{\epsilon_t}$$

# Renormalizing the Weight Distribution

- We need to ensure  $w_{t+1}(i)$  is a distribution (that sums up to 1)

$$\begin{aligned}w_{t+1}(i) &= \frac{\exp(-y_i H_{t+1}(\mathbf{x}_i))}{Z_{t+1}} && \text{(by definition)} \\&= \frac{1}{Z_{t+1}} \cdot \exp(-y_i [H_t(\mathbf{x}_i) + \alpha_{t+1} h_{t+1}(\mathbf{x}_i)]) \\&= \frac{1}{Z_{t+1}} \cdot (\exp(-y_i H_t(\mathbf{x}_i)) \cdot \exp(-y_i \alpha_{t+1} h_{t+1}(\mathbf{x}_i))) \\&= \frac{1}{Z_{t+1}} \cdot (w_t(i) \cdot \exp(-y_i \alpha_{t+1} h_{t+1}(\mathbf{x}_i)))\end{aligned}$$

# Comments

- This is just a theoretical explanation of the Adaboost algorithm. Exponential loss is not a good loss function anyway.
- Continuing to add weak classifiers after training error reaches 0 usually still increases the performance (due to the increase of the margin).
  - Similar phenomenon in Deep Learning.

## Old ML Conventional Wisdom

- Good prediction balances bias and variance.
- You should not perfectly fit your training data as some in-sample errors can reduce out-of-sample error.
- High capacity models don't generalize.
- Optimizing to high precision harms generalization.
- Nonconvex optimization is hard in machine learning.



## What we have *always* seen

- Interpolating your training data is fine.
- Training on your test set is fine.
- Making models huge doesn't hurt.
- Making models huge doesn't help much.

Source: Ben [Recht](#)

# Adaboost

---

**Algorithm 1:** Adaboost( $(\mathbf{X}, \mathbf{y}), T$ )

---

```
1  $w_0(i) = \frac{1}{n}, \forall i \in [n];$   
2 for  $t = 1$  to  $T$  do  
3    $h_t \leftarrow$  the weak classifier (from  $\mathcal{H}$ ) that minimizes weighted error  
   (see  $\epsilon_t$  below); /*  $h_t(\mathbf{x}) \in \{-1, 1\}$  */;  
4    $\epsilon_t \leftarrow \frac{1}{n} \cdot \sum_{i=1}^n w_{t-1}(i) \cdot \mathbf{1}[\![H(\mathbf{x}_i) \neq y_i]\!];$   
5    $\alpha_t \leftarrow \frac{1}{2} \ln \frac{1-\epsilon_t}{\epsilon_t};$   
6   for  $i = 1$  to  $n$  do  
7      $w_t(i) \leftarrow w_{t-1}(i) \cdot e^{-\alpha_t [y_i h_t(\mathbf{x}_i)]};$   
8   Renormalize  $w_{t+1}(i)$  to be a distribution;  
9 return  $H_T(\mathbf{x}) \stackrel{\text{def}}{=} \text{sgn} \left( \sum_{t=1}^T \alpha_t \cdot h_t(\mathbf{x}) \right);$ 
```

---

# References

- Scott Fortmann-Roe. Understanding the Bias-Variance Tradeoff. <http://scott.fortmann-roe.com/docs/BiasVariance.html>
- M. Magdon-Ismail. Approximation Versus Generalization. <http://www.cs.rpi.edu/~magdon/courses/LFD-Slides/SlidesLect07.pdf>
- Boris Babenko. Note: A Derivation of Discrete AdaBoost. [http://vision.ucsd.edu/~bbabenko/data/boosting\\_note.pdf](http://vision.ucsd.edu/~bbabenko/data/boosting_note.pdf).